

DATA PRIVACY **AND SECURITY**

Prof. Daniele Venturi

Master's Degree in Data Science
Sapienza University of Rome



CIS SAPIENZA

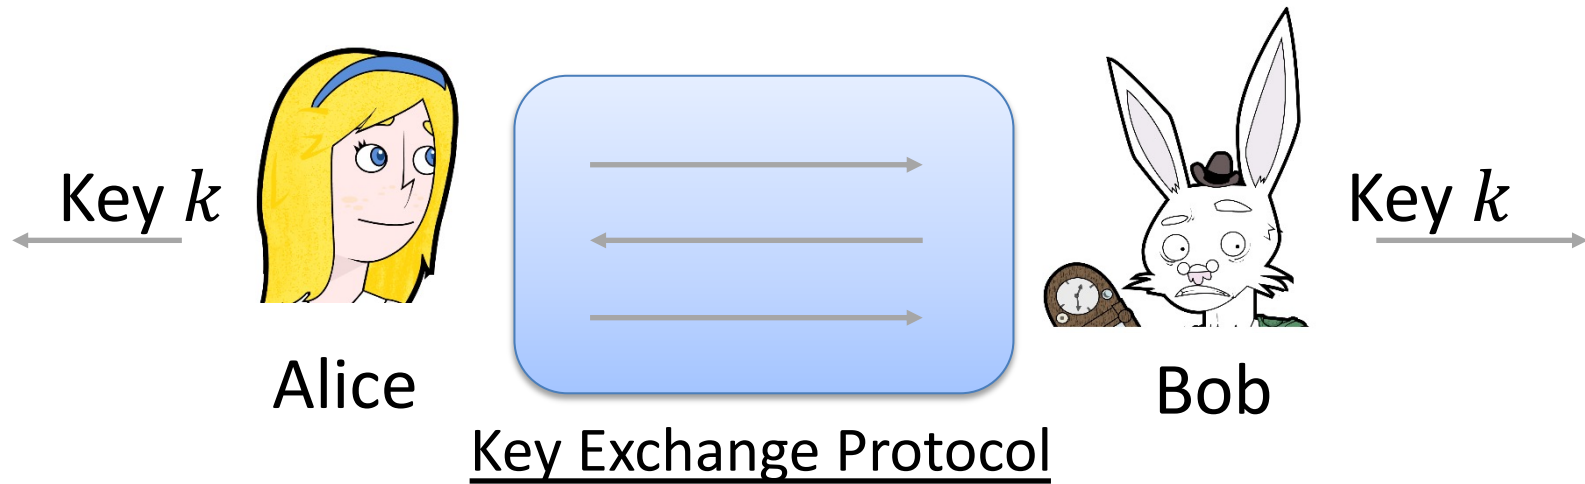
RESEARCH CENTER FOR CYBER INTELLIGENCE
AND INFORMATION SECURITY

CHAPTER 3:

Key Exchange




Key Exchange Protocols



- Allows to **agree** on a key over a **public** channel
 - KE bootstraps **secure communication**
 - KE constitutes the **link** between **symmetric** and **asymmetric** cryptography

Diffie-Hellman Key Exchange

$$k = (g^y)^x$$




$$x \leftarrow_{\$} \mathbb{Z}_q$$

$$g^x$$

$$g^y$$

$$k = (g^x)^y$$




$$y \leftarrow_{\$} \mathbb{Z}_q$$

- \mathbb{G} is a cyclic group of prime order q , with **generator** g
 - **Passive security** follows from DDH
 - E.g., \mathbb{G} is a subgroup of \mathbb{Z}_p^* where $q|p - 1$

Perfect Forward Secrecy

$$k = (g^y)^x$$




$$x \leftarrow_{\$} \mathbb{Z}_q$$

$$g^x$$

$$g^y$$


$$k = (g^x)^y$$



$$y \leftarrow_{\$} \mathbb{Z}_q$$

- Once the session keys are **destroyed** there is **no way** to recover them
 - Not even the owners (not even at gun point)

(Wo)Man-in-the-Middle Attack



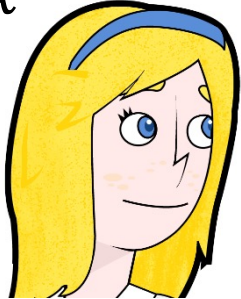
- Eve shares one secret key with **each party**
 - She can decrypt all subsequent communication
- Solution: **Authenticate** messages!
 - Master keys and session keys

Authenticated Key Exchange (AKE)

- Allow two parties to establish a **common secret** in an **authenticated** way
 - Parties should possess **previously established** authentication keys (master keys)
- **Secrecy**: The session key should be indistinguishable from a **random string**
- Additional properties:
 - **Mutual** authentication
 - **Consistency** (honest parties have a **consistent** view of **who** the peers to the session are)



First Attempt


$$X = g^x$$


$$x \leftarrow_{\$} \mathbb{Z}_q$$
$$sk_A, pk_B$$

$$K = Y^x$$

$$\xrightarrow{A, X, \mathbf{S}(sk_A, X)}$$

$$\xleftarrow{B, Y, \mathbf{S}(sk_B, Y)}$$




$$Y = g^y$$

$$y \leftarrow_{\$} \mathbb{Z}_q$$
$$sk_B, pk_A$$

$$K = X^y$$


- What if Eve ever finds an $(x, g^x, \mathbf{S}(sk_A, X))$?
 - Ephemeral leakage should not allow **long-term impersonation!**

Second Attempt




- View of the parties at the end of the protocol
 - A: Shared $K = g^{xy}$ with B
 - B: Shared $K = g^{xy}$ with A
 - **Looks fine, but...**

Identity-Misbinding Attack



- Wrong **identity binding**!
 - $A: K \Leftrightarrow B$, but $B: K \Leftrightarrow E$
- Eve doesn't know K , but Bob considers anything coming from Alice **as from Eve**

The ISO 9796 Defense



- Include the **peer identity** under the signature
 - Note that Eve cannot forge $\mathbf{S}(sk_B, X || Y || A)$
 - Avoids previous attack, and can be **proven secure**

Security Desiderata

- Intuitive (e.g., attacker capabilities, secrecy, ...)
- Reject **bad** protocols
- Accept **good** protocols
- Ensure security of **applications**
 - **Secure communication** in primis
 - **Composition** and usability
- We will overview the Canetti-Krawczyk (CK) model which is used to analyze many **real-world** KE protocols



Elements of the Definition

- A **two-party** protocol in a **multi-party** setting
- Multiple protocol executions run **concurrently**
 - Each run of a protocol at a party is called **session**
- Sessions are given **unique** names
 - (A, s_A) and (B, s_B) where B is the **intended peer**
 - The **session id** is (A, s_A, B, s_B)
 - Sessions with **corresponding** names like (A, s_A, B, s_B) and (A, s_A, B, s_B) are **matching**
 - At the end, a session outputs the **session id** and the **session key**

The Attacker

- We only assume **unauthenticated** channels
- The adversary
 - Monitors/controls/**modifies** traffic
 - Schedules sessions at will (**interleaving**)
 - May corrupt parties learning **long-term** secrets along with any state information and session keys
 - May issue learning queries for **short-term** information (e.g., session keys or state)
- A session is **exposed** if the owner is corrupted or the adversary issued learning query

The Security Definition

- Completed **matching** sessions output the same key (**correctness**)
- The attacker learns **nothing** about **unexposed** sessions
 - **Test session** chosen by the adversary
 - Attacker is given either the **honest** key or a **randomly generated** key and can't distinguish
 - Key **confirmation** can be added to the definition
- Note: Never use **session keys** as part of the KE protocol itself (e.g., TLS 1.2)

Sanity Checks


- The above definition is **simple** but **powerful**
 - **Impersonation**: If Eve can impersonate Bob **without** corrupting him, she **knows** a key for an **unexposed** session
 - Eve **can't** break one session given the key of **another** session
 - **Identity misbinding**: If Eve forces two (**non-matching**) sessions with outputs (A, B, K) and (B, E, K) , she can choose one to be the **test session** and use the other one to **expose** K



Authenticators

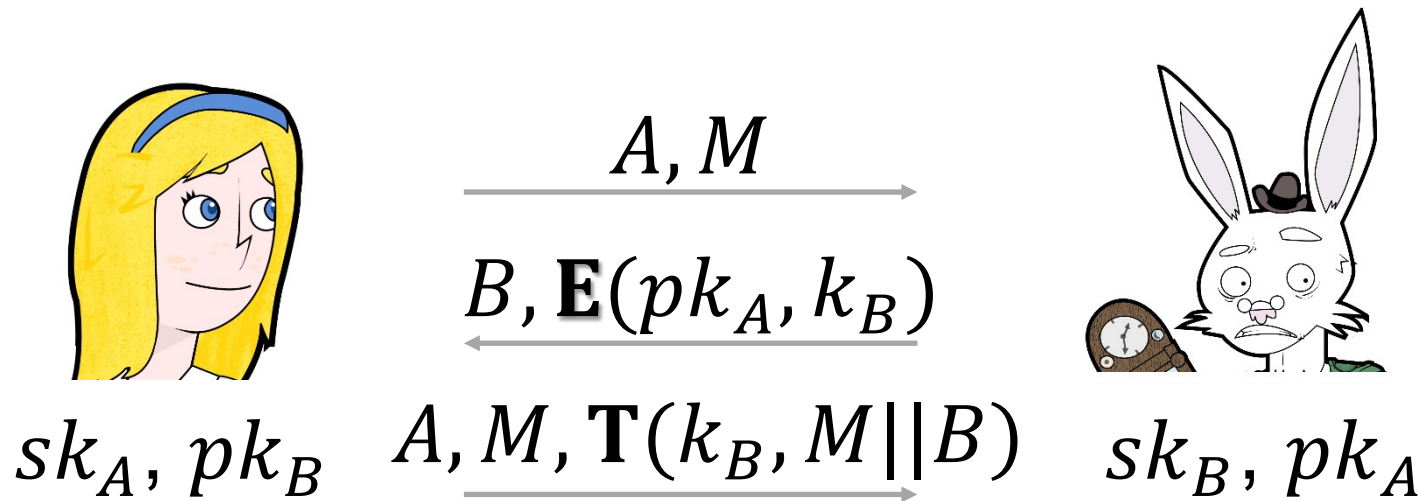
- Consider a much **weaker** attack model where a KE protocol **authenticated** channels
 - **Idealized** model with **passive** attacker
 - Still the attacker can do **everything else**
 - The DH protocol is **trivially secure** in this model
- Authenticators are protocol **compilers** that allow to **reduce** KE protocols secure in the **unauthenticated** channels model to ones in the **authenticated** channels model

Authenticators based on Signatures



- The nonce avoids **replay attacks**
- If Bob thinks that he **received** message M from Alice, then Alice **sent** M to Bob
 - One can show the above **implies** security of the ISO 9796 protocol in the CK model


Authenticators based on Encryption



- Alice is the **only** party that can **decrypt** the ciphertext sent by Bob
 - Under **randomly chosen** key k_B
- So Bob is convinced it received M from Alice
 - The first message can actually be **dropped** here

SKEME (IKEv1)

$$X = g^x$$



$$x \leftarrow_{\$} \mathbb{Z}_q$$
$$sk_A, pk_B$$

$$K = Y^x$$


$$\xrightarrow{A, \mathbf{E}(pk_B, k_A)}$$

$$\xleftarrow{B, Y, \mathbf{T}(k_A, Y || A), \mathbf{E}(pk_A, k_B)}$$

$$\xrightarrow{A, X, \mathbf{T}(k_B, X || B)}$$

$$k = \mathbf{H}(A || B || X || Y || K)$$

$$Y = g^y$$



$$y \leftarrow_{\$} \mathbb{Z}_q$$
$$sk_B, pk_A$$

$$K = X^y$$


- The keys k_A and k_B are **randomly** chosen
- Can be seen as applying the **encryption-based** authenticator on the classical DH protocol

On Identity Protection

- **Identity** protection
 - Hide identities from **passive/active** adversaries
- A **privacy** concern in many scenarios
 - Probing attacks in the internet
 - Location anonymity of roaming users
- The design of IPSec and IKE protocols is **heavily influenced** by the above concern
 - SKEME and SIGMA
 - Typically **only one** id is hidden in the presence of active adversaries




Why not ISO?



- **Unsuited** for identity protection
 - Bob **needs to know** Alice's identity and viceversa
 - Also, it leaves a **signed proof** of communication

SKEME with Encrypted IDs

$$X = g^x$$



$$x \leftarrow_{\$} \mathbb{Z}_q$$
$$sk_A, pk_B$$


$$K = Y^x$$

$$\xrightarrow{\mathbf{E}(pk_B, A || k_A)}$$

$$Y, \mathbf{T}(k_A, Y || A), \mathbf{E}(pk_A, B || k_B)$$

$$\xrightarrow{A, X, \mathbf{T}(k_B, X || B)}$$

$$Y = g^y$$




$$y \leftarrow_{\$} \mathbb{Z}_q$$
$$sk_B, pk_A$$

$$K = X^y$$

- The keys k_A and k_B are **randomly** chosen
- But Alice **needs to know** the public key of Bob **beforehand**

Alternative Solution: STS

$$X = g^x$$



$$x \leftarrow_{\$} \mathbb{Z}_q$$
$$sk_A, pk_B$$


$$K = Y^x$$

X

$$Y, \mathbf{E}(K, B || \mathbf{S}(sk_B, X || Y))$$

$$\mathbf{E}(K, A || \mathbf{S}(sk_A, X || Y))$$

$$Y = g^y$$




$$y \leftarrow_{\$} \mathbb{Z}_q$$
$$sk_B, pk_A$$

$$K = X^y$$

- Add a **proof of knowledge** of the secret key K
- **Insecure** if Eve can **register** pk_A as her key
 - At least in the variant where A is **in the clear**

STS using MACs

$$X = g^x$$



$$x \leftarrow_{\$} \mathbb{Z}_q$$

$$sk_A, pk_B$$


$$K = Y^x$$

$$X$$

$$Y, B, \sigma_B = \mathbf{S}(sk_B, X || Y), \mathbf{T}(K, \sigma_B)$$

$$A, \sigma_A = \mathbf{S}(sk_A, X || Y), \mathbf{T}(K, \sigma_A)$$

$$Y = g^y$$



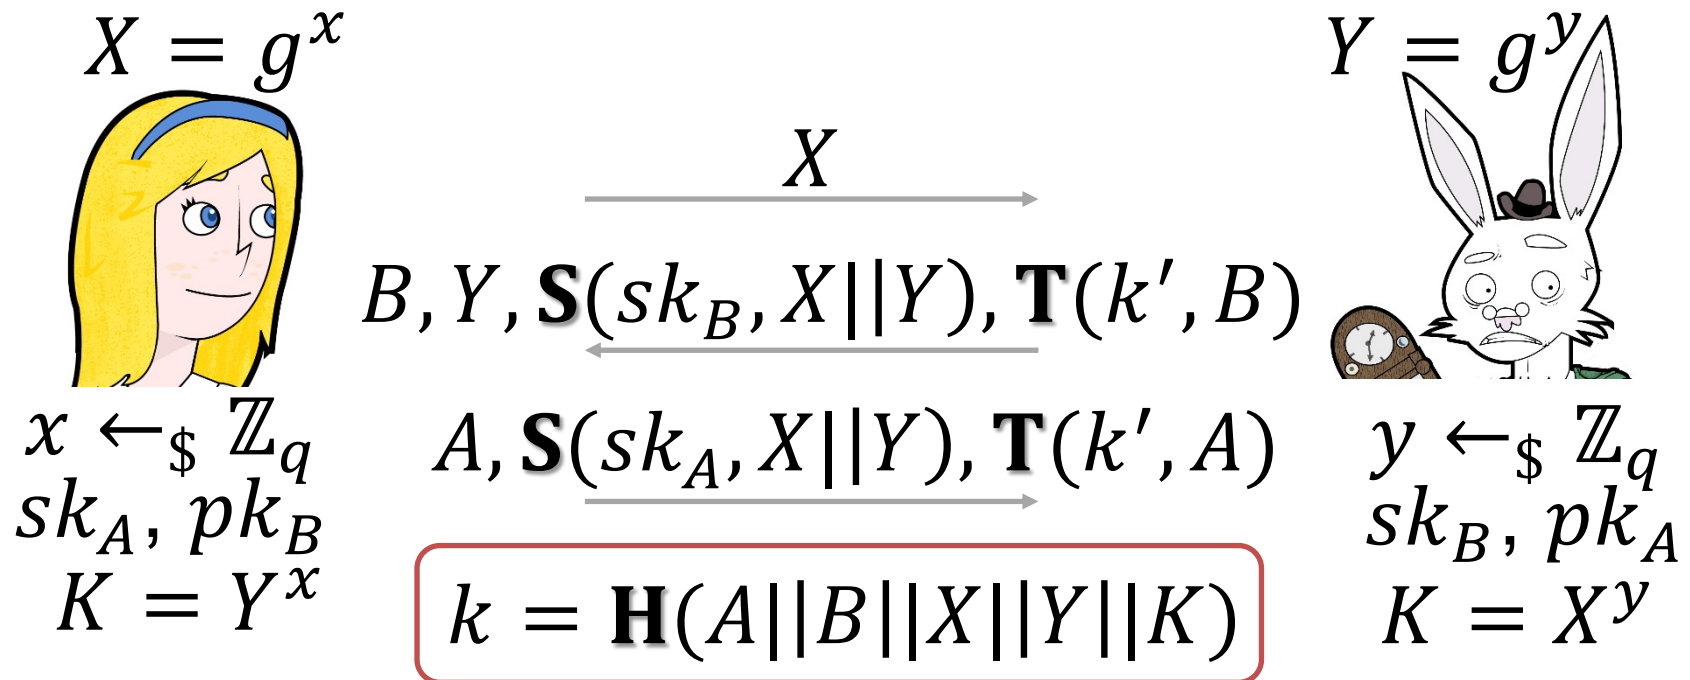
$$y \leftarrow_{\$} \mathbb{Z}_q$$

$$sk_B, pk_A$$

$$K = X^y$$

- MACs more suited to **prove knowledge** of K
- Yet, the **same attack** as before **still works**
 - We need to **bind** the **key** with the **peer ids**


SIGMA: Basic Version



- Instead of signing Alice's id (ISO), Bob **tags** its own identity with **another key** k'
 - The key k' is **derived** from K (as the session key k)


SIGMA-I: Protect Alice's ID (Initiator)

$$X = g^x$$



$$X$$


$$Y = g^y$$



$$Y, \mathbf{E}(k'', B || \mathbf{S}(sk_B, X || Y) || \mathbf{T}(k', B))$$


$$\mathbf{E}(k'', A || \mathbf{S}(sk_A, X || Y) || \mathbf{T}(k', A))$$

$$k = \mathbf{H}(A || B || X || Y || K)$$


$$\begin{aligned} x &\leftarrow_{\$} \mathbb{Z}_q \\ sk_A, pk_B \\ K &= Y^x \end{aligned}$$

$$\begin{aligned} y &\leftarrow_{\$} \mathbb{Z}_q \\ sk_B, pk_A \\ K &= X^y \end{aligned}$$

- **Encrypt** the identities of both Alice and Bob using **another key** k'' (still derived from k)
 - Bob's id is protected against **passive** attackers
 - Alice's id is protected against **active** attackers

SIGMA-R: Protect Bob's ID (Responder)


$$X = g^x$$



$$X$$

$$Y$$

$$Y = g^y$$



$$\mathbf{E}(k'', A || \mathbf{S}(sk_A, X || Y) || \mathbf{T}(k', A))$$

$$\mathbf{E}(k'', B || \mathbf{S}(sk_B, X || Y) || \mathbf{T}(k', B))$$

$$k = \mathbf{H}(A || B || X || Y || K)$$

$$\begin{aligned} x &\leftarrow_{\$} \mathbb{Z}_q \\ sk_A, pk_B \\ K &= Y^x \end{aligned}$$

$$\begin{aligned} y &\leftarrow_{\$} \mathbb{Z}_q \\ sk_B, pk_A \\ K &= X^y \end{aligned}$$

- Bob **does not** reveal his identity **before** checking who he is talking to
 - Bob's id is protected against **active** attackers
 - Alice's id is protected against **passive** attackers

Security of SIGMA


- The above description is **oversimplified** and glosses over a number of **details**
 - Additional information (context, negotiation, ...)
- Nevertheless, SIGMA can be **proved** secure in the CK model
 - But no **modular** proof using authenticators is currently known
- The protocol is used in IPSec as well as part of the new TLS 1.3 standard

AKE with Implicit Authentication

- Drawbacks of the ISO 9796 protocol
 - It requires to send **signatures** and **certificates**
- What is the **inherent cost** of **authentication**?
 - Communication complexity
 - Computation complexity
 - What security?
- **Implicit** authentication
 - No signatures or tags sent
 - **Ability to compute** session key → authentication


Only the certificates
are sent

Some Ideas




$$\xrightarrow{A = g^a, X = g^x}$$

$$\xleftarrow{B = g^b, Y = g^y}$$




- Many **insecure** attempts
 - $k = \mathbf{H}(g^{ab}, g^{xy})$: given a key for **one session** one can find a key for **another session**
 - $k = \mathbf{H}(g^{ab}, g^{xy}, g^x, g^y)$: knowing the key of Bob one can **impersonate** Alice to Bob
- **Want:** security unless (a, x) or (b, y) **leak**

MQV: The Basic Idea




$$A = g^a, X = g^x$$

$$B = g^b, Y = g^y$$



- **Idea:** Let $K = g^{(a+x)(b+y)}$
 - **Insecure:** Eve sends $X^* = g^{x^*}/A$; Bob sends Y , and thus $K = (BY)^{x^*} = AX^*$ which is the same as computed by Bob $(AX^*)^{b+y} = (BY)^{x^*}$
- **Avoid** the attack by letting $K = g^{(a+dx)(b+ey)}$
 - Values d, e s.t. Even **can't** control e, Y or d, X


Hashed MQV



a, x

$$A = g^a, X = g^x$$

$$B = g^b, Y = g^y$$



b, y

$$d = \mathbf{H}(X || \text{Bob})$$


$$e = \mathbf{H}(Y || \text{Alice})$$

$$K = (XA^d)^{y+be}$$

$$K = (YB^e)^{x+da}$$

- The **session key** is just $k = \mathbf{H}(K)$
 - Computing K requires $1 + 1/6$ **exponentiations**
- MQV: Let d be the **first half** bits of X and e be the **second half** bits of Y (but **insecure**)


Hashed MQV



a, x

$$A = g^a, X = g^x$$

$$B = g^b, Y = g^y$$



b, y

$$d = \mathbf{H}(X||Bob)$$

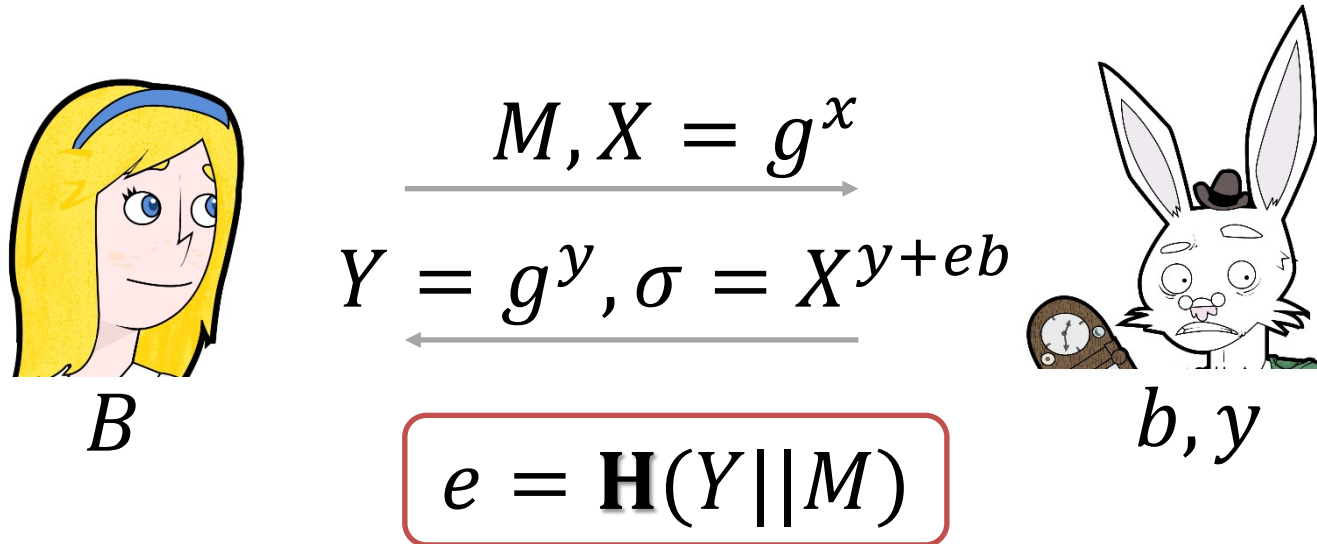
$$e = \mathbf{H}(Y||Alice)$$

$$K = (YB^e)^{x+da}$$

$$K = (XA^d)^{y+be}$$


- No signatures **exchanged**
 - But we can think of $(YB^e)^{x+da}$ (resp. $(XA^d)^{y+be}$) as a **signature** of Alice on $X||Bob$ (resp. $Y||Alice$)
 - **Same** signature by **different** parties on **different** messages

XCR Signatures



- Bob is the **signer** with public key $B = g^b$
 - Alice sends a **message** M and a **challenge** $X = g^x$
 - Alice **accepts** iff $(YB^e)^x = \sigma$
- Alice is a **designated verifier**

Dual XCR Signatures




B, a, x

$\xrightarrow{M_A, X \cdot A^d}$

$\xleftarrow{M_B, Y \cdot B^e}$

$$\begin{aligned} d &= \mathbf{H}(X || M_B) \\ e &= \mathbf{H}(Y || M_A) \end{aligned}$$



A, b, y

- Alice and Bob act as **simultaneous** signers
 - Bob (Alice) generates an **XCR signature** on **challenge** $X \cdot A^d$ ($Y \cdot B^e$) and **message** M_A (M_B)
 - **Same** signature $\sigma = (XA^d)^{y+eb} = (YB^e)^{x+da}$

Security of HMQV


- One can show that HMQV is **secure** in the CK model (assuming **H** is a **random oracle**)
 - Reduce **security** of HMQV to **unforgeability** of Dual XCR signatures
 - Reduce **unforgeability** of Dual XCR signatures to **unforgeability** of XCR signatures
 - Reduce **unforgeability** of XCR signatures to the **CDH assumption** in the **random oracle model**
- The protocol is **standardized** by ANSI/ISO and IEEE, and also used by the NSA



Key Derivation Functions (KDFs)

- A KDF turns an **imperfect** source of randomness into one or more **random keys**
 - **Imperfect:** Not uniform
- In practice one just uses **random oracles**
 - As in $k = \mathbf{H}(g^{xy})$
 - Repeated extraction as $\mathbf{H}(g^{xy} || A) || \mathbf{H}(g^{xy} || B) \dots$
- However, **no \mathbf{H}** can be a **random oracle**
 - **Length extension attack:** Given $\mathbf{H}(g^{xy} || A)$ can compute $\mathbf{H}(g^{xy} || B)$ if A is a **prefix** of B

Extract-then-Expand



- The value s is a **salt** that is **public** but **random**
 - This is usually also **short**
- The value K is the starting **key material**
- Extract function: a **randomness extractor**
- Expand function: typically a **PRF**


Instantiations in Practice

- There are **statistically-secure** extractors
 - But in **practice** those would require **large seeds** and yield quite **large entropy loss**
- **Alternative:** Use a PRF for **both** extraction and expansion
 - Difficulty: the seed is **public** (but the input is **not**)
 - There are **examples** of PRFs that **do not work**
- Luckily, the above works using **practical** PRFs
 - In particular, with the **standard** HMAC

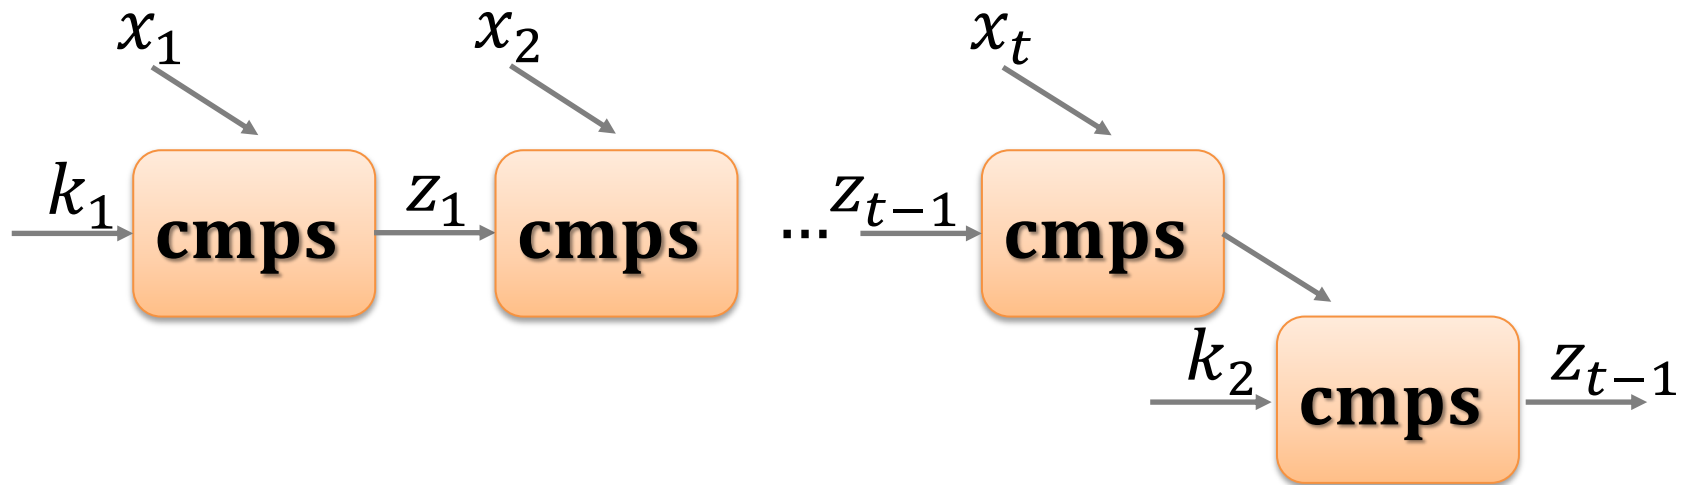


Keyed Merkle-Damgaard

- Let **cmps** be a **compression function** outputting 160 bits out of 512 bits
- The **keyed** Merkle-Damgaard construction uses the seed s as **initial vector**




NMAC: PRF Mode for Merkle-Damgaard



- **Theorem:** $\text{NMAC}(k_1 || k_2, \cdot)$ is a PRF assuming **cmps** is a PRF
- HMAC is identical, but k_1, k_2 are **derived** from the **same key** k

Extract-than-Expand



- Expand function:

$$k_{i+1} = \mathbf{HMAC}(k_{\text{prf}}, k_{i+1} || \text{info} || i)$$

- This is HMAC as a PRF in **feedback mode**
- Heavily **standardized** (e.g., TLS 1.3, Whatsapp)
 - And also **provably secure**

Applications of HKDF

- IPsec:
 - $k = \mathbf{HKDF}(\text{nonces}, g^{xy})$ where the **nonces** are part of the protocol and used as **salt**
 - In case the nonces are **public** the analysis requires that **HKDF** is an **extractor**
 - In case the nonces are **secret** (SKEME) the analysis requires that **HKDF** is a **PRF**
- TLS 1.3 with shared key \hat{k} (**resumption**):
 - $k = \mathbf{HKDF}(\hat{k}, g^{xy})$
 - **HKDF** as an **extractor/PRF** if \hat{k} is **revealed/secret**

Password-Authenticated Key Exchange

- **Authenticated** key exchange still requires a **public-key infrastructure**
- Alternative: Rely on a **shared password**
- The **standardization** of PAKE took several years starting back in 1982
- Today, PAKE is used in many **use cases**
 - TLS 1.3 (**pre-shared** key variant)
 - iCloud
 - RFID authentication



Passwords

- A **password** is a string of symbols belonging to a finite alphabet
 - Equivalently a bitstring
 - Needs to be **stored securely**
- Typical applications:
 - Derive a cryptographic key
 - Password-based authentication



Attacks on Passwords

- Guessing **always** possible (brute force)
 - **Online**: Trial & error
 - **Offline**: Dictionary attacks
- Sniffing from networks or theft from server
- Software attacks (trojan horse programs)
- Social engineering (phishing)
- Shoulder surfing



Online Password Guessing

- Always possible
 - Servers are always online
- Requires **interaction** with server
 - Limit number of **failed attempts**
 - Limit guessing rate
- Guessing rate
 - Attempt failure counter (but can't block user account)
 - Increasing answer delay after each failed attempt

Offline Password Guessing

- Can't be **detected**
- Attacker may choose **amount of resources**
- Complexity of guessing can be controlled by careful **password selection**
 - Given value $y = f(\pi, z)$, where f, z are public, a guessing attempt π' means to check $y = f(\pi', z)$



Passwords Entropy

- Let X be a random variable outputting symbols from an alphabet $\mathcal{A} = \{a_1, \dots, a_n\}$
- Denote by p_i the probability associated to a_i
- **Average information** in bit/symbol

$$H(X) = - \sum_{i=1}^n p_i \log p_i$$

- Maximum entropy for uniform distribution
 $H(U) = \log n$

ASCII Passwords

- Consider 7 bit ASCII: 95 **printable chars**
 - 0-31 are control chars
 - 127 is a special char
- For uniform passwords, with $n = 95$ we have $H(U) = \log 95 = 6.57$ bit/char
 - 128 bits of security correspond to random password of roughly 20 chars
- Situation **gets worse** if only upper/lower chars and numbers are used
 - $H(U) = \log 62 = 5.95$ bit/char



Passphrases

- More often users choose **passphrases**
- Let $p(\vec{x})$ be the probability of ℓ consecutive chars $\vec{x} = (x_1, \dots, x_\ell) \in \mathcal{A}^\ell$

- Now

$$H(X) = \lim_{\ell \rightarrow \infty} \frac{-\sum_{\vec{x} \in \mathcal{A}^\ell} p(\vec{x}) \log p(\vec{x})}{\ell}$$

$H_\ell(X)$

- Italian language: $H_3(X) \approx 3.15$ bit/char;
 $H_5(X) \approx 2.22$ bit/char; $H_6(X) \approx 1.87$
bit/char

Users Choose Poor Passwords

- Study at Purdue University (1992)

Length	Number	Fraction of Total
1	55	0.4
2	87	0.6
3	212	2
4	449	3
5	1260	9
6	3035	22
7	2917	21
8	5772	42%

- Among 69 million Yahoo! Passwords, 1.1% of users pick **same password**


Password Selection

- Computer generated and **refreshed**
 - Difficult to remember!
- System process periodically tries **guessing** user passwords
 - CPU intensive
 - Memory intensive for big dictionaries
 - Users might get annoyed
- Check user password **as entered**
 - Simple guidance to select acceptable passwords



Bloom Filters (1/2)

- Tradeoff between accuracy and time/memory to check passwords **belong to dictionary** \mathcal{D}
- Let \mathbf{H}_i be k hash functions yielding values in $[0, N - 1]$ for $N = 2^s$ and T a table of N bits
- Let $y_i = \mathbf{H}_i(w)$, $\forall w \in \mathcal{D}$ and set $T[y_i] = 1$
- Given π , reject it iff $T[\mathbf{H}_i(\pi)] = 1, \forall i \in [k]$




Bloom Filters (2/2)

- If $\pi \in \mathcal{D}$, it is **always rejected**
- If $\pi \notin \mathcal{D}$, it **might be rejected** (false positive)
 - Let $q = \Pr[T[j] = 0 : j \in [0, N - 1]] = \Pr[\mathbf{H}_i(w) \neq j : \forall i \in [k], w \in \mathcal{D}]$
- False positive rate:

$$p = (1 - q)^k = (1 - (1 - 1/N)^{kD})^k \approx (kD2^{-s})^k$$
- Optimal values for fixed false positive rate:
$$k \approx -\log_2 p ; N \approx -1.44 \cdot D \cdot \log_2 p$$

Password based Encryption

PKCS#5 Standard




$\mathbf{E}(\pi, M)$:

$\text{salt} \leftarrow_{\$} \{0,1\}^{128}$

$K = \mathbf{H}^c(\pi || \text{salt})$


$C = K \oplus M$

Output (salt, C)



e.g., $c = 10000$

Salt and Stretching



- Hash chain **slows down** attacks by factor of c
- Salt defeats rainbow tables and provides **separation** between users

Typically assumed to be **trivial** for the adversary

Step 1:

$$M_1 = \mathbf{H}^c(\pi_1 || \text{salt}) \oplus C$$

$$M_2 = \mathbf{H}^c(\pi_2 || \text{salt}) \oplus C$$

$$M_3 = \mathbf{H}^c(\pi_3 || \text{salt}) \oplus C$$

...

Step 2:

$$M_1 = \text{as7e657q622!|a1}$$

$$M_2 = \text{mnas237@##saw}$$


$$M_3 = \text{sometext}$$

...



C, salt

Honey Encryption



Seems **indistinguishable** to the adversary



Step 1:

$$M_1 = \mathbf{H}^c(\pi_1 || \text{salt}) \oplus C$$

$$M_2 = \mathbf{H}^c(\pi_2 || \text{salt}) \oplus C$$

$$M_3 = \mathbf{H}^c(\pi_3 || \text{salt}) \oplus C$$

...

Step 2:

$$M_1 = 01010000111000$$

$$M_2 = 01111100011000$$

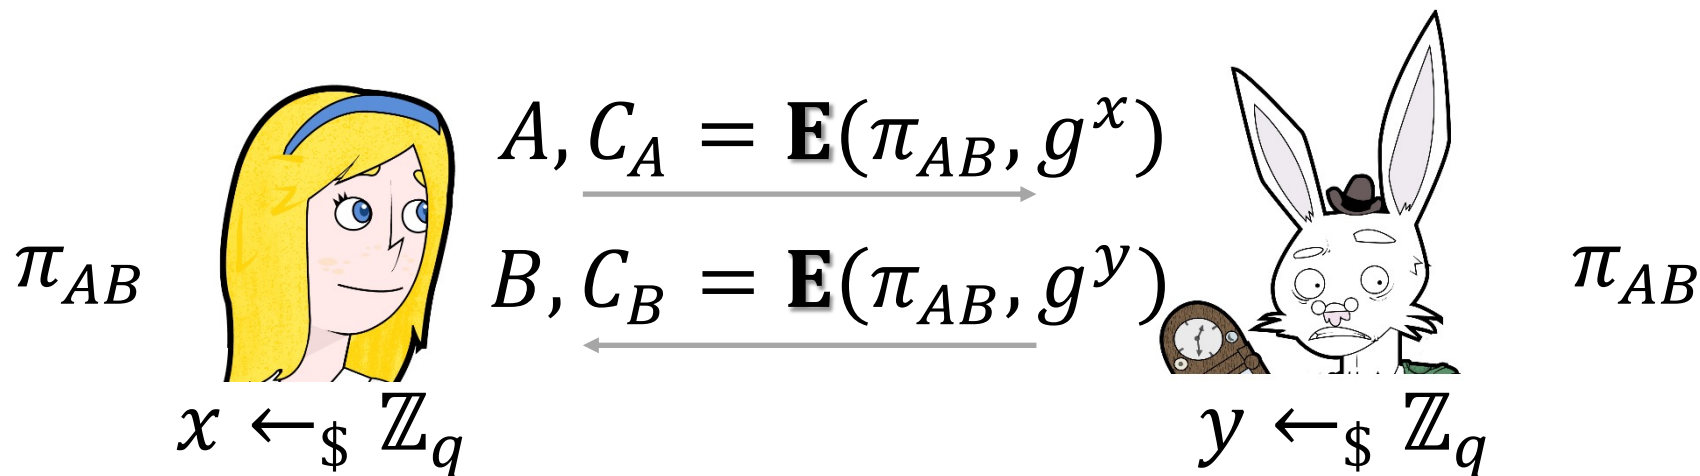
$$M_3 = 11001111000101$$

...

Encrypted Key Exchange (EKE)

$$k = \mathbf{D}(\pi_{AB}, C_B)^x$$

$$k = \mathbf{D}(\pi_{AB}, C_A)^y$$




- Instantiation:
 - $\mathbf{E}(\pi, M) = \text{ideal cipher}$
 - Hash protocol transcript with a **random oracle**

Transport Layer Security (TLS)

- Goal: Establish a **secure channel**
 - **Key exchange**: Yields keys for confidentiality/authenticity
 - **Record layer**: Use keys to secure communication
 - Authentication (usually on server side)
- Used in tons of applications
 - Amazon, ebay, e-commerce
 - Email
 - Google



The Client-Sever Scenario



- What actually happens:
 - You type amazon.it in your browser
 - TLS connection with Amazon is negotiated
 - You get to https:// for **secure browsing**
 - You **authenticate** to Amazon on a **secure link**

History of TLS

- Started out as Secure Socket Layer (SSL)
 - Developed by Netscape around 1995
 - Goal: Secure communication over Internet
- Changed to TLS in 1999
 - Secure communication (HTTPS)
 - ... but also FTP, secure emailing, etc.
 - **Heavily standardized**
- Many implementations
 - OpenSSL, BoringSSL, s2n (TLS by Amazon)

SSL/TLS Versions

- SSL 1.0: Never released
 - Too **insecure** for release
- SSL 2.0: Released in February 1995
 - But contained a number of **security flaws**
- SSL 3.0: Released in 1996
- TLS 1.1: Protection against CBC-mode attacks
- TLS 1.2: Move from MD5 to SHA-1 (2008)
 - However, first attacks on MD5 **already in 2005**
- TLS 1.3: August 2018; completely **revised**

Attacks on TLS

- Renegotiation attack on SSL 3.0
 - **Ideal patch:** Kill renegotiation
 - **Real patch:** include previous session history
- Version rollback attacks
 - **Ideal patch:** Kill backward compatibility
 - **Real patch:** ??? (not a realistic attack)
- BEAST: Browser exploits of CBC vulnerabilities
 - **Ideal patch:** Kill CBC mode
 - **Real patch:** Discourage CBC mode

Attacks on TLS (cont'd)

- Lucky 13: Exploit padding problems
 - Ideal patch: Kill CBC mode
 - Real patch: encouraged RC4 or use AES-GCM
- POODLE: Downgrade to SSL 3.0
 - Ideal patch: Kill backward compatibility
 - Real patch: ???



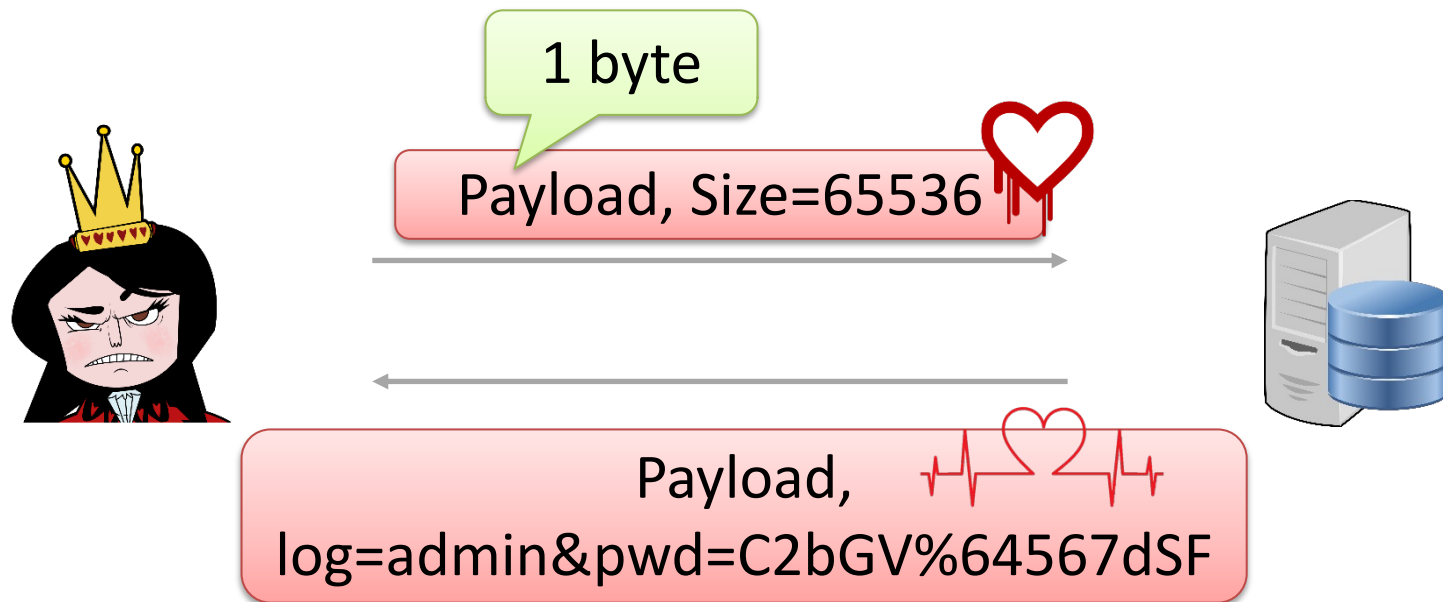
Even More Attacks

- RC4 attacks: RC4 output is biased
 - Ideal patch: Kill RC4
 - Real patch: RFC 7465 prohibits RC4, but
 - 30% of TLS traffic still uses RC4
 - 75% of sites allow RC4 negotiation
- Heartbleed, 3Shake, FREAK, Logjam
- ...




Heartbleed

- Attack on OpenSSL based on **HeartBeats**
 - HeartBeat requests keep a TLS connection alive
 - HeartBeat contains a payload along with its size



TLS 1.3: (EC)DHE



handshake key

**ClientHello
ClientKeyShare**



**ServerHello
ServerKeyShare**



**ServerConfiguration
ServerCertificate
ServerCertificateVerify
ServerFinished**



**ClientCertificate
ClientCertificateVerify
ClientFinished**




channel key



handshake key

channel key

TLS 1.3: Crypto Details




$$N_C \leftarrow \{0,1\}^{256}$$
$$x \leftarrow \mathbb{Z}_q$$

handshake key

$$\mathbf{KDF}(g^{xy}, CH, \dots, SKS)$$

channel key

$$\mathbf{KDF}(g^{xy}, CH, \dots, CF)$$



$$N_S \leftarrow \{0,1\}^{256}$$
$$y \leftarrow \mathbb{Z}_q$$

handshake key

$$\mathbf{KDF}(g^{xy}, CH, \dots, SKS)$$


$$\sigma = \mathbf{S}(sk_S, CH, \dots, SCert)$$

$$\tau = \mathbf{T}(k_{SF}, CH, \dots, SKS)$$

channel key

$$\mathbf{KDF}(g^{xy}, CH, \dots, CF)$$

TLS 1.3: Pre-Shared Key Variant



preshared key

Externally or
from session
resumption

ClientHello
ClientKeyShare
early_data
psk_ke_modes
psk_shared_key →

ServerHello
ServerKeyShare
psk_shared_key
encrypted_extensions
ServerFinished ←

⋮




preshared key

Zero Round-Trip Time

- TLS 1.3 requires a few messages before a key is established
- ORTT is an alternative to the PSK variant
- The client starts the protocol and immediately delivers data
 - This is achieved using a semi-static server key
 - This key is available for short time periods
 - ORTT was first invented by Google in order to reduce the latency

ORTT: QUIC



semi-static

server key g^s

ephemeral key e, g^e

$$k_1 = \mathbf{KDF}(g^{es})$$

$$k_2 = \mathbf{KDF}(g^{et})$$



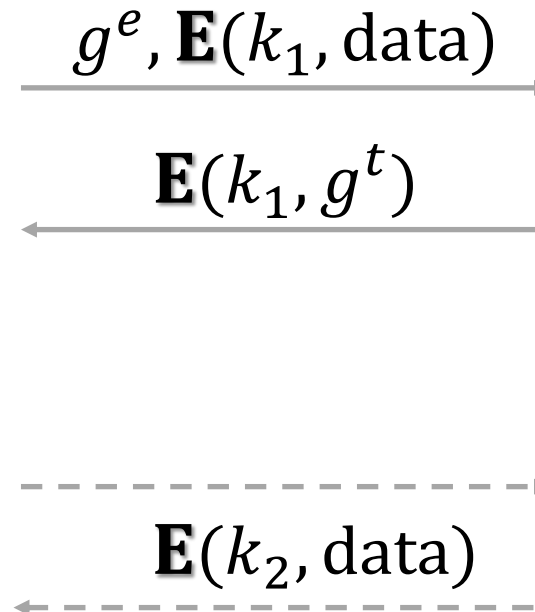
semi-static

server key s


$$k_1 = \mathbf{KDF}(g^{es})$$

ephemeral key t, g^t

$$k_2 = \mathbf{KDF}(g^{et})$$



Replay Attacks on QUIC



semi-static

server key g^s

ephemeral key e, g^e

$$k_1 = \mathbf{KDF}(g^{es})$$



$g^e, \mathbf{E}(k_1, \text{data})$

$g^e, \mathbf{E}(k_1, \text{data})$



semi-static

server key s

$$k_1 = \mathbf{KDF}(g^{es})$$

Only way out:
Store previously
received values